

# Santiago\_Gabriela

February 6, 2022

```
[18]: !pip install missingno
```

```
Requirement already satisfied: missingno in
/Users/gabby/opt/anaconda3/lib/python3.9/site-packages (0.5.0)
Requirement already satisfied: matplotlib in
/Users/gabby/opt/anaconda3/lib/python3.9/site-packages (from missingno) (3.4.3)
Requirement already satisfied: numpy in
/Users/gabby/opt/anaconda3/lib/python3.9/site-packages (from missingno) (1.20.3)
Requirement already satisfied: seaborn in
/Users/gabby/opt/anaconda3/lib/python3.9/site-packages (from missingno) (0.11.2)
Requirement already satisfied: scipy in
/Users/gabby/opt/anaconda3/lib/python3.9/site-packages (from missingno) (1.7.1)
Requirement already satisfied: cycler>=0.10 in
/Users/gabby/opt/anaconda3/lib/python3.9/site-packages (from
matplotlib->missingno) (0.10.0)
Requirement already satisfied: kiwisolver>=1.0.1 in
/Users/gabby/opt/anaconda3/lib/python3.9/site-packages (from
matplotlib->missingno) (1.3.1)
Requirement already satisfied: pillow>=6.2.0 in
/Users/gabby/opt/anaconda3/lib/python3.9/site-packages (from
matplotlib->missingno) (8.4.0)
Requirement already satisfied: python-dateutil>=2.7 in
/Users/gabby/opt/anaconda3/lib/python3.9/site-packages (from
matplotlib->missingno) (2.8.2)
Requirement already satisfied: pyparsing>=2.2.1 in
/Users/gabby/opt/anaconda3/lib/python3.9/site-packages (from
matplotlib->missingno) (3.0.4)
Requirement already satisfied: six in
/Users/gabby/opt/anaconda3/lib/python3.9/site-packages (from
cycler>=0.10->matplotlib->missingno) (1.16.0)
Requirement already satisfied: pandas>=0.23 in
/Users/gabby/opt/anaconda3/lib/python3.9/site-packages (from seaborn->missingno)
(1.3.4)
Requirement already satisfied: pytz>=2017.3 in
/Users/gabby/opt/anaconda3/lib/python3.9/site-packages (from
pandas>=0.23->seaborn->missingno) (2021.3)
```

```
[1]: #import libraries: pandas, numpy, matplotlib (set %matplotlib inline),
      ↪matplotlib's pyplot, seaborn, missingno, scipy's stats, sklearn
      #This jupyter notebook is prepared by "Gabriela Santiago".
      %matplotlib inline
      import pandas as pd
      import matplotlib.pyplot as plt
      import numpy as np
      import seaborn as sns
      import scipy.stats as st
      import sklearn
      import missingno as msno
```

```
[2]: #import the data to a dataframe and show how many rows and columns does it have
      df = pd.read_csv('hrdata.csv')
```

```
[3]: print("Rows: ", len(df))
      print("Columns: ", len(df.columns))
```

Rows: 21287  
Columns: 18

```
[9]: #call the describe method of dataframe to see some summary statistics of the
      ↪numerical columns
      df.describe()
```

```
[9]:
```

	Unnamed: 0	rec_num	enrollee_id	city_development_index \
count	21287.000000	21287.000000	21287.000000	21287.000000
mean	10643.000000	10644.000000	16873.983652	0.828462
std	6145.171926	6145.171926	9612.131237	0.123537
min	0.000000	1.000000	1.000000	0.448000
25%	5321.500000	5322.500000	8554.500000	0.739000
50%	10643.000000	10644.000000	16967.000000	0.903000
75%	15964.500000	15965.500000	25161.500000	0.920000
max	21286.000000	21287.000000	33380.000000	0.949000

	training_hours	target	city_development_matrices
count	21287.000000	19158.000000	21287.000000
mean	65.328510	0.249348	8.284615
std	60.075201	0.432647	1.235365
min	1.000000	0.000000	4.480000
25%	23.000000	0.000000	7.390000
50%	47.000000	0.000000	9.030000
75%	88.000000	0.000000	9.200000
max	336.000000	1.000000	9.490000

```
[10]: #Explain in words if you find any column's statistics interesting and good to
      ↪know
```

*#I think it's good to know the training hours' statistics to know if an  
 ↳ employee has average or above average training time*

```
[11]: #Show the top 5 rows and last 5 rows of the data frame
df.head()
```

```
[11]: Unnamed: 0  rec_num  enrollee_id      city  city_development_index  gender \
0           0         1        8949  city_103                0.920  Male
1           1         2       29725  city_40                0.776  Male
2           2         3       11561  city_21                0.624  NaN
3           3         4       33241  city_115               0.789  NaN
4           4         5         666  city_162               0.767  Male

      relevent_experience  enrolled_university  education_level \
0  Has relevent experience      no_enrollment      Graduate
1  No relevent experience      no_enrollment      Graduate
2  No relevent experience  Full time course      Graduate
3  No relevent experience                NaN      Graduate
4  Has relevent experience      no_enrollment      Masters

      major_discipline  experience  company_size  company_type  last_new_job \
0           STEM          >20         NaN         NaN          1
1           STEM          15       50-99      Pvt Ltd          >4
2           STEM           5         NaN         NaN      never
3  Business Degree          <1         NaN      Pvt Ltd      never
4           STEM          >20       50-99  Funded Startup          4

      training_hours  target  state  city_development_matrices
0           36      1.0    CA                9.20
1           47      0.0    CA                7.76
2           83      0.0    CA                6.24
3           52      1.0    CA                7.89
4            8      0.0    CA                7.67
```

```
[12]: df.tail()
```

```
[12]: Unnamed: 0  rec_num  enrollee_id      city  city_development_index \
21282      21282      21283        1289  city_103                0.920
21283      21283      21284         195  city_136                0.897
21284      21284      21285       31762  city_100                0.887
21285      21285      21286        7873  city_102                0.804
21286      21286      21287       12215  city_102                0.804

      gender  relevent_experience  enrolled_university  education_level \
21282  Male  No relevent experience      no_enrollment      Graduate
21283  Male  Has relevent experience      no_enrollment      Masters
21284  Male  No relevent experience      no_enrollment  Primary School
```

21285	Male	Has relevent experience	Full time course	High School
21286	Male	Has relevent experience	no_enrollment	Masters

	major_discipline	experience	company_size	company_type	last_new_job \
21282	Humanities	16	NaN	Public Sector	4
21283	STEM	18	NaN	NaN	2
21284	NaN	3	NaN	Pvt Ltd	never
21285	NaN	7	100-500	Public Sector	1
21286	STEM	15	10000+	Pvt Ltd	2

	training_hours	target	state	city_development_matrices
21282	15	NaN	CA	9.20
21283	30	NaN	CA	8.97
21284	18	NaN	CA	8.87
21285	84	NaN	CA	8.04
21286	11	NaN	CA	8.04

```
[13]: #List all the numerical columns
numerical = df.select_dtypes(include=[np.number])
numerical.columns
```

```
[13]: Index(['Unnamed: 0', 'rec_num', 'enrollee_id', 'city_development_index',
        'training_hours', 'target', 'city_development_matrices'],
        dtype='object')
```

```
[14]: #List all the categorical columns
categorical = df.select_dtypes(include=[object])
categorical.columns
```

```
[14]: Index(['city', 'gender', 'relevent_experience', 'enrolled_university',
        'education_level', 'major_discipline', 'experience', 'company_size',
        'company_type', 'last_new_job', 'state'],
        dtype='object')
```

```
[15]: #Show a list with column wise count of missing values and display the list in
      ↪count wise descending order
nulls = df.isnull().sum().to_frame('nulls')
nulls.sort_values("nulls", inplace = True, ascending = False)
for index, row in nulls.iterrows():
    print(index, row[0])
```

```
company_type 6774
company_size 6560
gender 5016
major_discipline 3125
target 2129
education_level 512
last_new_job 463
```


```
enrolled_university 417
experience 70
state 0
training_hours 0
Unnamed: 0 0
rec_num 0
relevent_experience 0
city_development_index 0
city 0
enrollee_id 0
city_development_matrices 0
```

```
[16]: #Show a list with column wise percentage of missing values and display the list
      ↪ in percentage wise descending order
percentage = df.isnull().mean()*100
percentage = percentage.to_frame("nulls")
percentage.sort_values("nulls", inplace = True, ascending = False)
for index, row in percentage.iterrows():
    print(index, row[0])
```

```
company_type 31.822238925165593
company_size 30.816930520975244
gender 23.563677361770093
major_discipline 14.680321322873116
target 10.001409310846995
education_level 2.405223845539531
last_new_job 2.1750364071968806
enrolled_university 1.9589420773241883
experience 0.3288391976323578
state 0.0
training_hours 0.0
Unnamed: 0 0.0
rec_num 0.0
relevent_experience 0.0
city_development_index 0.0
city 0.0
enrollee_id 0.0
city_development_matrices 0.0
```



```
[17]: #Display a bar plot to visualize only the columns with missing values and their
      ↪ count. The plot should display from less missing value columns in the left
      ↪ and then more missing value columns to the right side of the plot
missing = df.isnull().sum()
missing = missing[missing>0]
missing.sort_values(inplace = True)
plt.figure(figsize = (10, 7))
missing.plot.bar()
```

[17]: <AxesSubplot:>



```
[31]: #Use missingno's bar plot, matrix plot with 200 sample, and heatmap
      msno.matrix(df.sample(200))
      msno.heatmap(df)
```


[31]: <AxesSubplot:>



[32]: *#Interpret any interesting information you found in the heatmap and any one plot*  
*#It's noticeable that there is a relatively strong correlation between the*  
*↪ 'company\_size' values and the "company\_type" category*


*#It's also noticeable that those two very categories happen to be the columns  
→with the most missingno variables, this may have biased the heatmap*

```
[20]: #Use seaborn bar plot for the categorical feature to see different values and  
      ↪count  
      #I used countplots instead as the bar plots required a y value  
      sns.countplot(x = 'city', data = categorical)  
      sns.set(rc={'figure.figsize':(1, 13)})
```







```
[83]: sns.countplot(x = 'gender', data = categorical)
sns.set(rc={'figure.figsize':(10, 13)})
```



```
[35]: sns.countplot(x = 'relevent_experience', data = categorical)
sns.set(rc={'figure.figsize':(10, 13)})
```




```
[36]: sns.countplot(x = 'enrolled_university', data = categorical)
sns.set(rc={'figure.figsize':(10, 13)})
```




```
[37]: sns.countplot(x = 'education_level', data = categorical)
```

```
[37]: <AxesSubplot:xlabel='education_level', ylabel='count'>
```




```
[38]: sns.countplot(x = 'major_discipline', data = categorical)
```

```
[38]: <AxesSubplot:xlabel='major_discipline', ylabel='count'>
```




```
[39]: sns.countplot(x = 'experience', data = categorical)
```

```
[39]: <AxesSubplot:xlabel='experience', ylabel='count'>
```




```
[40]: sns.countplot(x = 'company_size', data = categorical)
```

```
[40]: <AxesSubplot:xlabel='company_size', ylabel='count'>
```



```
[41]: sns.countplot(x = 'company_type', data = categorical)
```


```
[41]: <AxesSubplot:xlabel='company_type', ylabel='count'>
```



```
[42]: sns.countplot(x = 'last_new_job', data = categorical)
```


```
[42]: <AxesSubplot:xlabel='last_new_job', ylabel='count'>
```






```
[43]: sns.countplot(x = 'state', data = categorical)
```


```
[43]: <AxesSubplot:xlabel='state', ylabel='count'>
```




```
[44]: #Use seaborn countplot for the categorical feature against the values of the  
      →target  
sns.countplot(x = 'city', hue = df['target'], data = categorical)  
sns.set(rc={'figure.figsize':(10, 130)})
```



```
[21]: sns.countplot(x = 'gender', hue = df['target'], data = categorical)
sns.set(rc={'figure.figsize':(10, 13)})
```




```
[46]: sns.countplot(x = 'relevent_experience', hue = df['target'], data = categorical)
sns.set(rc={'figure.figsize':(10, 13)})
```




```
[47]: sns.countplot(x = 'enrolled_university', hue = df['target'], data = categorical)
```

```
[47]: <AxesSubplot:xlabel='enrolled_university', ylabel='count'>
```




```
[48]: sns.countplot(x = 'education_level', hue = df['target'], data = categorical)
```

```
[48]: <AxesSubplot:xlabel='education_level', ylabel='count'>
```



```
[49]: sns.countplot(x = 'major_discipline', hue = df['target'], data = categorical)
```


```
[49]: <AxesSubplot:xlabel='major_discipline', ylabel='count'>
```



```
[50]: sns.countplot(x = 'experience', hue = df['target'], data = categorical)
```


```
[50]: <AxesSubplot:xlabel='experience', ylabel='count'>
```






```
[51]: sns.countplot(x = 'company_size', hue = df['target'], data = categorical)
```

```
[51]: <AxesSubplot:xlabel='company_size', ylabel='count'>
```




```
[52]: sns.countplot(x = 'company_type', hue = df['target'], data = categorical)
```

```
[52]: <AxesSubplot:xlabel='company_type', ylabel='count'>
```




```
[53]: sns.countplot(x = 'last_new_job', hue = df['target'], data = categorical)
```

```
[53]: <AxesSubplot:xlabel='last_new_job', ylabel='count'>
```



```
[54]: sns.countplot(x = 'state', hue = df['target'], data = categorical)
```

```
[54]: <AxesSubplot:xlabel='state', ylabel='count'>
```




[55] : *#Interpret any interesting information and any information that might help you*  
*→to make any decision on combining, removing, or adding features based on*  
*→that, or any resampling maybe needed.*  
*#Since everyone in the data is from California I don't believe we need to keep*  
*→that column at all as it doesn't contribute to any learning about the data*  
*→or its correlation to the target value.*

```
#It also looks like STEM graduates not enrolled in a university make up for
↳most of the people who wish to stay with the company, which could be useful
↳for creating a predictive model
#City 103 has the most people willing to stay for the company, since every
↳other city has a miniscule amount of people, I would change the data to a
↳binary input of "living in city 103" and "not living in city 103"
```

```
[56]: #For each numerical features, perform the following:
      #Plot their distributions using histogram
      print(numerical.columns)
      plt.hist(df['rec_num'])
```

```
Index(['Unnamed: 0', 'rec_num', 'enrollee_id', 'city_development_index',
      'training_hours', 'target', 'city_development_matrices'],
      dtype='object')
```


```
[56]: (array([2129., 2129., 2128., 2129., 2128., 2129., 2129., 2128., 2129.,
      2129.]),
      array([1.00000e+00, 2.12960e+03, 4.25820e+03, 6.38680e+03, 8.51540e+03,
      1.06440e+04, 1.27726e+04, 1.49012e+04, 1.70298e+04, 1.91584e+04,
      2.12870e+04]),
      <BarContainer object of 10 artists>)
```



```
[57]: plt.hist(df['enrollee_id'])
```

```
[57]: (array([2013., 2123., 2087., 2130., 2118., 2167., 2144., 2165., 2135.,  
          2205.]),  
      array([1.00000e+00, 3.33890e+03, 6.67680e+03, 1.00147e+04, 1.33526e+04,  
          1.66905e+04, 2.00284e+04, 2.33663e+04, 2.67042e+04, 3.00421e+04,
```


```
3.33800e+04]],  
<BarContainer object of 10 artists>)
```



```
[58]: plt.hist(df['city_development_index'], density = True)  
plt.figure(figsize = (10, 10))
```

```
[58]: <Figure size 720x720 with 0 Axes>
```






<Figure size 720x720 with 0 Axes>


```
[59]: plt.hist(df['training_hours'], density = True)
```

```
[59]: (array([0.01131796, 0.00806743, 0.0045154 , 0.00243579, 0.00163929,  
            0.00062402, 0.0004852 , 0.00025802, 0.00021595, 0.00029168]),  
      array([ 1. , 34.5, 68. , 101.5, 135. , 168.5, 202. , 235.5, 269. ,  
            302.5, 336. ]),  
      <BarContainer object of 10 artists>)
```




```
[60]: plt.hist(df['target'], density = True)
```

```
[60]: (array([7.50652469, 0.          , 0.          , 0.          , 0.          ,  
          0.          , 0.          , 0.          , 0.          , 2.49347531]),  
array([0. , 0.1, 0.2, 0.3, 0.4, 0.5, 0.6, 0.7, 0.8, 0.9, 1. ]),  
<BarContainer object of 10 artists>)
```



```
[61]: plt.hist(df['city_development_matrices'], density = True, bins = 2)
```


```
[61]: (array([0.09551058, 0.30369102]),  
      array([4.48 , 6.985, 9.49 ]),  
      <BarContainer object of 2 artists>)
```



```
[62]: #Plot the distribution using seaborn distplot
sns.distplot(df['rec_num'])
```

```
/Users/gabby/opt/anaconda3/lib/python3.9/site-
packages/seaborn/distributions.py:2619: FutureWarning: `distplot` is a
deprecated function and will be removed in a future version. Please adapt your
code to use either `displot` (a figure-level function with similar flexibility)
or `histplot` (an axes-level function for histograms).
  warnings.warn(msg, FutureWarning)
```

```
[62]: <AxesSubplot:xlabel='rec_num', ylabel='Density'>
```




```
[63]: sns.distplot(df['enrollee_id'])
```

/Users/gabby/opt/anaconda3/lib/python3.9/site-packages/seaborn/distributions.py:2619: FutureWarning: `distplot` is a deprecated function and will be removed in a future version. Please adapt your code to use either `displot` (a figure-level function with similar flexibility)

```
or `histplot` (an axes-level function for histograms).  
warnings.warn(msg, FutureWarning)
```

```
[63]: <AxesSubplot:xlabel='enrollee_id', ylabel='Density'>
```




```
[64]: sns.distplot(df['city_development_index'])
```

```
/Users/gabby/opt/anaconda3/lib/python3.9/site-  
packages/seaborn/distributions.py:2619: FutureWarning: `distplot` is a  
deprecated function and will be removed in a future version. Please adapt your  
code to use either `displot` (a figure-level function with similar flexibility)  
or `histplot` (an axes-level function for histograms).  
warnings.warn(msg, FutureWarning)
```

```
[64]: <AxesSubplot:xlabel='city_development_index', ylabel='Density'>
```






```
[65]: sns.distplot(df['training_hours'])
```

```
/Users/gabby/opt/anaconda3/lib/python3.9/site-  
packages/seaborn/distributions.py:2619: FutureWarning: `distplot` is a  
deprecated function and will be removed in a future version. Please adapt your  
code to use either `displot` (a figure-level function with similar flexibility)
```

```
or `histplot` (an axes-level function for histograms).  
warnings.warn(msg, FutureWarning)
```


```
[65]: <AxesSubplot:xlabel='training_hours', ylabel='Density'>
```



```
[66]: sns.distplot(df['target'])
```

```
/Users/gabby/opt/anaconda3/lib/python3.9/site-  
packages/seaborn/distributions.py:2619: FutureWarning: `distplot` is a  
deprecated function and will be removed in a future version. Please adapt your  
code to use either `displot` (a figure-level function with similar flexibility)  
or `histplot` (an axes-level function for histograms).  
warnings.warn(msg, FutureWarning)
```


```
[66]: <AxesSubplot:xlabel='target', ylabel='Density'>
```



```
[67]: sns.distplot(df['city_development_matrices'])
```

```
/Users/gabby/opt/anaconda3/lib/python3.9/site-  
packages/seaborn/distributions.py:2619: FutureWarning: `distplot` is a  
deprecated function and will be removed in a future version. Please adapt your  
code to use either `displot` (a figure-level function with similar flexibility)  
or `histplot` (an axes-level function for histograms).  
warnings.warn(msg, FutureWarning)
```

```
[67]: <AxesSubplot:xlabel='city_development_matrices', ylabel='Density'>
```




[68]: *#Interpret any interesting information*

```
#These plots don't work well with things like IDs because there is one for each  
↪ person, but it's good for making observations with values like  
↪ 'training_hours' where you can observe that the majority of people have  
↪ between 0 and around 70 hours of training under their belt, and there are  
↪ some with as many as 300 training hours
```

```
[69]: #For the numerical attributes, use heatmap to show the correlation  
correlation = numerical.corr()  
sns.heatmap(correlation)
```

```
[69]: <AxesSubplot:>
```




```
[70]: #If you find any interesting short list of columns, create another heatmap with
      ↪ them and show the correlations inside the heatmap as well
      #Nothing interesting I hadn't already explained from the above heatmap
```

```
[71]: #Show scatter plots between columns to show the relationships with the target
      plt.scatter(data = df, x = 'training_hours', y = 'target')
```

```
[71]: <matplotlib.collections.PathCollection at 0x7fd25238df70>
```






[72]: *#As you can see, a scatterplot in these situations wouldn't be very useful for*  
*↪ analysis as the values of target are either 1 or 0*

[73]: *#Interpret and explain any finding and next course of action from there*

```
#Since the city_development_index and the city_development_matrix have no  
↪correlation to the target, there's no real reason to use it in a predictive  
↪model. I'd remove them both
```

```
[74]: #Use boxplot or any other strategies to find outliers  
plt.boxplot(data = df, x = 'rec_num')
```


```
[74]: {'whiskers': [<matplotlib.lines.Line2D at 0x7fd25240df40>,  
                  <matplotlib.lines.Line2D at 0x7fd25241f310>],  
       'caps': [<matplotlib.lines.Line2D at 0x7fd25241f6a0>,  
                <matplotlib.lines.Line2D at 0x7fd25241fa30>],  
       'boxes': [<matplotlib.lines.Line2D at 0x7fd25240dbb0>],  
       'medians': [<matplotlib.lines.Line2D at 0x7fd25241fdc0>],  
       'fliers': [<matplotlib.lines.Line2D at 0x7fd25242c190>],  
       'means': []}
```



```
[75]: plt.boxplot(data = df, x = 'enrollee_id')
```


```
[75]: {'whiskers': [<matplotlib.lines.Line2D at 0x7fd269b27dc0>,  
                  <matplotlib.lines.Line2D at 0x7fd269b9a160>],  
       'caps': [<matplotlib.lines.Line2D at 0x7fd269b9afd0>,  
                <matplotlib.lines.Line2D at 0x7fd269b9a8b0>],
```

```
'boxes': [<matplotlib.lines.Line2D at 0x7fd26bf8d250>],  
'medians': [<matplotlib.lines.Line2D at 0x7fd26a0c26a0>],  
'fliers': [<matplotlib.lines.Line2D at 0x7fd26994a070>],  
'means': []}
```




```
[76]: plt.boxplot(data = df, x = 'city_development_index')
```

```
[76]: {'whiskers': [<matplotlib.lines.Line2D at 0x7fd25293d2e0>,  
    <matplotlib.lines.Line2D at 0x7fd25293d670>],  
    'caps': [<matplotlib.lines.Line2D at 0x7fd25293da00>,  
    <matplotlib.lines.Line2D at 0x7fd25293dd90>],  
    'boxes': [<matplotlib.lines.Line2D at 0x7fd252933f10>],  
    'medians': [<matplotlib.lines.Line2D at 0x7fd252949160>],  
    'fliers': [<matplotlib.lines.Line2D at 0x7fd2529494f0>],  
    'means': []}
```



```
[77]: plt.boxplot(data = df, x = 'training_hours')
```

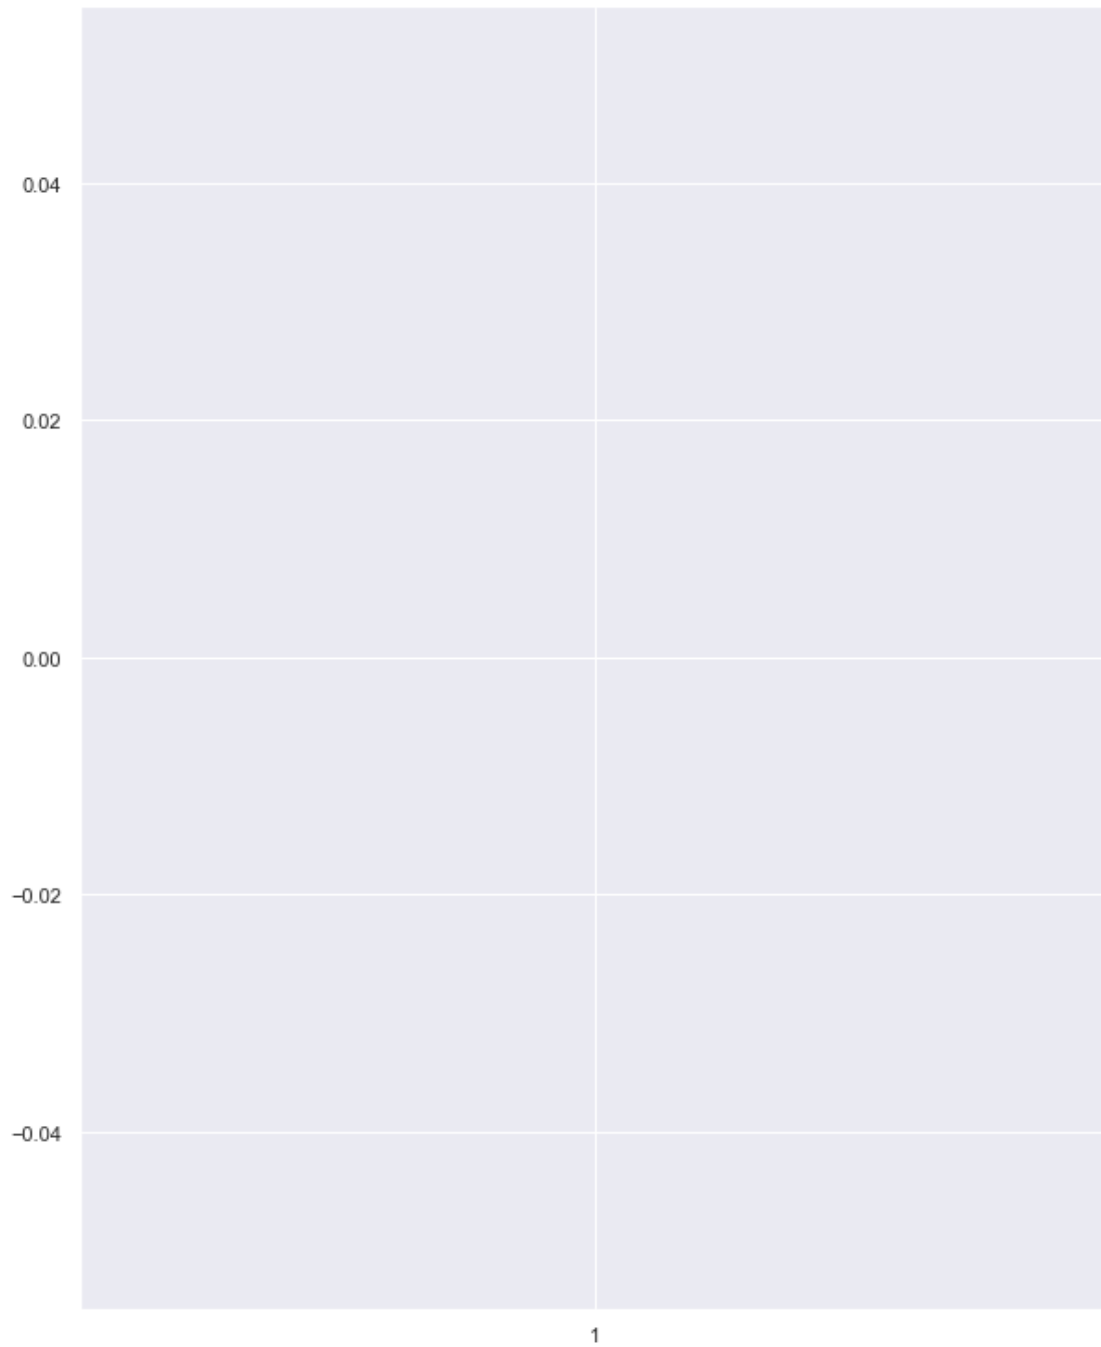
```
[77]: {'whiskers': [<matplotlib.lines.Line2D at 0x7fd253495580>,  
                 <matplotlib.lines.Line2D at 0x7fd253495910>],  
       'caps': [<matplotlib.lines.Line2D at 0x7fd253495ca0>,  
               <matplotlib.lines.Line2D at 0x7fd2534a0070>],  
       'boxes': [<matplotlib.lines.Line2D at 0x7fd2534951f0>],  
       'medians': [<matplotlib.lines.Line2D at 0x7fd2534a0400>],  
       'fliers': [<matplotlib.lines.Line2D at 0x7fd2534a0790>],  
       'means': []}
```



```
[78]: plt.boxplot(data = df, x = 'target')
```

```
[78]: {'whiskers': [<matplotlib.lines.Line2D at 0x7fd2536a58e0>,  
                 <matplotlib.lines.Line2D at 0x7fd2536a5c70>],  
      'caps': [<matplotlib.lines.Line2D at 0x7fd2536b3040>],
```


```
<matplotlib.lines.Line2D at 0x7fd2536b33d0>],  
'boxes': [<matplotlib.lines.Line2D at 0x7fd2536a5550>],  
'medians': [<matplotlib.lines.Line2D at 0x7fd2536b3760>],  
'fliers': [<matplotlib.lines.Line2D at 0x7fd2536b3af0>],  
'means': []}
```





```
[79]: plt.boxplot(data = df, x = 'city_development_matrices')
```

```
[79]: {'whiskers': [<matplotlib.lines.Line2D at 0x7fd253a78b20>,
<matplotlib.lines.Line2D at 0x7fd253a78eb0>],
'caps': [<matplotlib.lines.Line2D at 0x7fd253a86280>,
<matplotlib.lines.Line2D at 0x7fd253a86610>],
'boxes': [<matplotlib.lines.Line2D at 0x7fd253a78790>],
'medians': [<matplotlib.lines.Line2D at 0x7fd253a869a0>],
'fliers': [<matplotlib.lines.Line2D at 0x7fd253a86d30>],
'means': []}
```



```
[80]: #Lots of outliars in the 'training_hours' column, only 1 in the  
      ↪ 'city_development_matrices' column, and 1 in the 'city_development_index'
```

[81]: #What are the different values of experience, can you categorize them in to 0,  
 ↳1, and 2?  
 #The years of experience in this data set are values that range from less than  
 ↳one year of experience to more than 20 years.  
 #It wouldn't be too difficult to batch them into groups of 3, like 0 = years  
 ↳ranging from <1 to 7, 1 = years from 8 to 14, and 2 = years from 15 to 20+

[82]: #Finally after all the above EDA, summarize your finding, next course of action,  
 ↳such as we may need to transform distribution because of right skew etc,  
 ↳need to remove a particular columns for any reasons, remove records for any  
 ↳reasons, need to rebalance data and what are the rebalancing options (if  
 ↳needed), and any other finding  
 #After the above EDA, values like rec\_num, enrollee\_id, and state are all  
 ↳columns I'd remove from the data as they the first two are unique for each  
 ↳person, thus no correlation can be found from that. Also, the state column  
 ↳consists of only those from california, so it's not useful for predicting or  
 ↳analysis at all.  
 #I would also consider not using any data that is missing a target value as it  
 ↳wouldn't be very useful in predicting they they'd stay with the company or  
 ↳not.  
 #It appears that men with relevant experience, are not currently enrolled in a  
 ↳university, are university graduates from the STEM field with more than 20  
 ↳years of experience are only a few categories that make up the majority of  
 ↳those trained that wish to stay with the company. I believe the data  
 ↳provided would be sufficient for creating and training a predictive model,  
 ↳so long as the columns specified above are removed, as the model created  
 ↳might have a bias against anyone in a state other than california.